

Lab 6: reliable communication

191220138 杨飞洋

实验目的

学习路由器的工作原理，并通过代码实现其中的一部分。

具体功能为：接收并回复ICMP包，并对一些异常情况做出处理。

背景知识

interface接口

```
class switchyard.lib.interface.Interface(name, ethaddr, ipaddr=None, netmask=None, ifnum=None, iftype=<InterfaceType.Unknown: 1>)
```

ethaddr

Get the Ethernet address associated with the interface

ifnum

Get the interface number (integer) associated with the interface

iftype

Get the type of the interface as a value from the InterfaceType enumeration.

ipaddr

Get the IPv4 address associated with the interface

ipinterface

Returns the address assigned to this interface as an IPInterface object. (see documentation for the built-in ipaddress module).

name

Get the name of the interface

netmask

Get the IPv4 subnet mask associated with the interface

packet接口

get_header(hdrclass, returnval=None)

Return the first header object that is of class hdrclass, or None if the header class isn't found.

get_header_by_name(hdrname)

Return the header object that has the given (string) header class name. Returns None if no such header exists.

get_header_index(hdrclass, startidx=0)

Return the first index of the header class hdrclass starting at startidx (default=0), or -1 if the header class isn't found in the list of headers.

has_header(hdrclass)

Return True if the packet has a header of the given hdrclass, False otherwise.

UDP接口

UDP (user datagram protocol) header

`class switchyard.lib.packet.UDP(*kwargs)`

[\[source\]](#)

The UDP header contains just source and destination port fields.

`dst`

`length`

`src`

To construct a packet that includes an UDP header as well as some application data, the same pattern of packet construction can be followed:

```
>>> p = Ethernet() + IPv4(protocol=IPProtocol.UDP) + UDP()
>>> p[UDP].src = 4444
>>> p[UDP].dst = 5555
>>> p += b'These are some application data bytes'
>>> print (p)
Ethernet 00:00:00:00:00:00->00:00:00:00:00:00 IP | IPv4 0.0.0.0->0.0.0.0 UDP | UDP 4444->5555 | RawPacketContents (37 bytes) b'These are ...
>>>
```

python byte操作

[资料链接](#)

实现逻辑

middlebox

`middlebox` 类似一个路由器，只不过它只有两个端口，从这个端口进入，就从另一个端口发出，不用转发等操作，非常方便。

我们可以从 `start_mininet.py` 中读到关于端口的信息，以及 `blastee`、`blaster` 的 `mac`、`ip` 地址。

```
def setup_addressing(net):
    '''
    * reset_macs call sets the MAC address of the nodes in the network
    * blaster and blastee has a single port, hence the MAC address ends with :01
    * middlebox has two ports, MAC address ends with :01 and :02 respectively,
    that are connected to the blaster and blastee.
    '''
    reset_macs(net, 'blaster', '10:00:00:00:00:{:02x}'.format(0))
    reset_macs(net, 'blastee', '20:00:00:00:00:{:02x}'.format(0))
    reset_macs(net, 'middlebox', '40:00:00:00:00:{:02x}'.format(0))
    '''
    * set_ip_pair call assigns IP addresses of the interfaces
    * convention is same as MAC address
    * middlebox has two IP addresses: 192.168.100.2 and 192.168.200.2 - connected
    to blaster and blastee respectively
    '''
    set_ip_pair(net,
        'blaster', 'middlebox', '192.168.100.1/30', '192.168.100.2/30')
    set_ip_pair(net,
        'blastee', 'middlebox', '192.168.200.1/30', '192.168.200.2/30')
    set_route(net, 'blaster', '192.168.200.0/24', '192.168.100.2')
    set_route(net, 'blastee', '192.168.100.0/24', '192.168.200.2')
```

从上可以看到 `middlebox` 和 `blastee` 相连的端口 `mac` 地址是 `40:00:00:00:00:02`，`middlebox` 和 `blaster` 相连的端口 `mac` 地址是 `40:00:00:00:00:01`，`blaster` 的 `mac` 地址是 `10:00:00:00:00:01`，`blastee` 的 `mac` 地址是 `20:00:00:00:00:01`。

需要处理一个特殊情况就是当 blaster 发给 blastee 时，可能会丢包，这是一个概率事件，概率是 `dropRate="0.19"`

那可以用 `randint(1,100)` 函数来生成一个1到100的随机数来判断是否丢包，如果丢包了，就打印一下序列号，如果没丢包，就配置一下mac地址，从另一个端口发出去。

```
if fromIface == "middlebox-eth0":
    _drop = randint(1,100)
    if _drop >= 100 * self.dropRate:
        eth_idx = packet.get_header_index(Ethernet)
        packet[eth_idx].src = '40:00:00:00:00:02'
        packet[eth_idx].dst = '20:00:00:00:00:01'
        self.net.send_packet("middlebox-eth1", packet)
    else:
        sequence = packet[3].to_bytes()[:4]
        seq = int.from_bytes(sequence, 'big')
        print(seq, " packet dropped")
```

如果是从 blastee 发给 blaster 的则不用考虑丢包，配置一下地址直接从另一个端口发就行了：

```
elif fromIface == "middlebox-eth1":
    eth_idx = packet.get_header_index(Ethernet)
    packet[eth_idx].src = '40:00:00:00:00:01'
    packet[eth_idx].dst = '10:00:00:00:00:01'
    self.net.send_packet("middlebox-eth0", packet)
```

这就完成了 middlebox 的工作了。

关于python中 `int` 和 `byte` 之间的相互转换，并且用大端方式，在背景知识

blastee

blastee 需要回发一个ack包，ack包的格式为：

```
<----- Switchyard headers -----> <----- Your packet header(raw bytes) ----->
<-- Payload in raw bytes --->

-----
-----

|  ETH Hdr |  IP Hdr  |  UDP Hdr  |          Sequence number(32 bits)          |
|          | Payload  (8 bytes) |          |          |
-----
-----
```

首先肯定是设置包头，目的mac地址是 middlebox 与之相连的端口mac，目的ip是 blaster 的ip地址，ip层协议为UDP

blastee 接收到的来自 blaster 的数据包的格式如下：

```

----- Switchyard headers -----> <----- Your packet header(raw bytes) ----->
<-- Payload in raw bytes ---
-----
-----
|  ETH Hdr  |  IP Hdr  |  UDP Hdr  |  Sequence number(32 bits) | Length(16 bits) |
|  variable length payload  |
-----
-----

```

首先从packet获取序列号，是数据的前四个字节，长度是数据区的第5、6个字节，数据区的剩余就是数据。

序列号不变，直接放进ack包中。

ack包不用字节来指示长度。

当读取的packet包的 `length < 8` 时，说明数据不够8个字节，需要补0来对齐。反之，当读取的packet包的 `length >= 8` 时，说明此时数据字节数肯定够，则读 [6:14] 这前8个字节就行了。

这里需要python中int和byte相互转化，并且用大端方式，可在[背景知识: python-byte操作](#)可以找到教程。

实现代码如下：

```
ack = Ethernet() + IPv4(protocol=IPProtocol.UDP) + UDP()
ack[0].src = '20:00:00:00:00:01'
ack[0].dst = '40:00:00:00:00:02'
ack[1].src = '192.168.200.1'
ack[1].dst = '192.168.100.1'
ack[1].ttl = 10
sequence = packet[3].to_bytes()[:4]
#payload = packet[3].to_bytes[6:14]
ack += sequence
length = int.from_bytes(packet[3].to_bytes()[4:6], byteorder = 'big')
if length < 8:
    ack += packet[3].to_bytes()[6:]
    ack += (0).to_bytes(8 - length, byteorder = "big")
else:
    ack += packet[3].to_bytes()[6:14]
self.net.send_packet(fromIface,ack)
```

blaster

主要实现一个滑动窗口功能。

滑动窗口的原理在[手册](#)中十分清晰，这里就不过多赘述了，主要来谈如何实现的。

在 `blaster` 类中定义一个队列 `queue` 来存放。

```
self.queue = []
for i in range(0,self.num+1):
    self.queue.append(mission())
```

定义一个 `mission` 类来存放各个序列号对应的包的状态

```
class mission:
    def __init__(self) -> None:
        self.is_acked = 0
        self.is_sent = 0
```

`is_acked` 指示是否收到了ack包, `is_sent` 指示有没有进行重发, 来防止多余的重发。

首先来看下 `__init__()` 函数:

除了必要的参数之外, 还初始化了这个队列, 并且初始化了 `timer` 这个用于指示重发的时间变量, 还有一系列的需要在最后打印的数值。

还有很重要的 `rhs, lhs` 来指示队列的左右指针。

```
self.net = net
# TODO: store the parameters
self.blasterIp = blasterIp
self.num = int(num)
self.length = int(length)
self.sw = int(senderwindow)
self.timeout = int(timeout)
self.recv_to = int(recvTimeout)
self.lhs = self.rhs = 1
self.begintime = self.endtime = -1
self.ack_num = 0 //收到ack的数目
self.timer = time.time()
self.queue = []
self.re_transmit_num = 0 // 重传次数
self.num_timeout = 0 //超时次数
self.output_byte = 0 //总长度
self.good_put = 0 //有效长度
self.last_seq = -1 //上一次重传的seq
for i in range(0, self.num+1):
    self.queue.append(mission())
```

接下来看 `blaster` 的工作逻辑:

`handle_packet`

如果收到了包, 则执行 `handle_packet()` 函数, 在这个实验中, 很显然收到的包一定是从 `blastee` 发来的ack包, 那只需要识别出序列号, 将对应序列号的 `is_acked` 置为1, 并且如果序列号是 `self.lhs`, 将 `self.lhs` 后移, 并在此时更新 `blaster` 的时间。

但也是有情况是 `self.lhs` 不会移动, 比如:

window : 3 4 5 6 7

滑动窗口是 3 4 5 6 7, 但是在3的包丢包了, 收到了来自4 5 6 7的ack包, 此时就不会移动ack, 直到重发再收到3的ack包。

如果这个序列号的ack包此前没有被收到, 则表示这是第一次收到, `goodput+=length` 且 `ack_num++`, 如果 `ack_num == num`, 表示收到了所有的ack包, 此时就可以设置 `endtime` 了。

代码如下:

```

sequence = packet[3].to_bytes()[:4]
seq = int.from_bytes(sequence, 'big')
print("ack ", seq)
if self.queue[seq].is_acked == 0:
    self.good_put += self.length
    self.queue[seq].is_acked = 1
    self.ack_num += 1
if self.ack_num == self.num:
    self.endtime = time.time()
i = 1
while i < self.num+1 and self.queue[i].is_acked == 1:
    i += 1
if i > self.lhs:
    self.lhs = i
    if self.lhs > self.rhs:
        self.rhs = self.lhs
    self.timer = time.time()

```

handle_no_packet

当没有收到ack包时，就执行 `handle_no_packet()` 函数。

在这个函数中，我会优先考虑是否超时了，如果超时了，代表序列号为 `self.lhs` 的包没有收到ack，即在 `middlebox` 上丢包了，此时会重发这个包。

`re_transmit_num++` 即重发数目自增。

比较上次重发的seq即 `last_seq`，如果不一样则 `num_timeout++`。

用 `process_pkt()` 函数来包装一下这个包，从端口发出，同时 `output_byte += length`。

```

if time.time() - self.timer > self.timeout:
    self.re_transmit_num += 1
    if self.last_seq != self.lhs:
        self.last_seq = self.lhs
        self.num_timeout += 1
    self.process_pkt(pkt, self.lhs)
    self.output_byte += self.length
    self.net.send_packet("blaster-eth0", pkt)
    print(self.lhs, " timeout resend")

```

说下 `process_pkt()` 函数的实现，因为之后还会复用这个函数。

首先来看下数据包的结构：

```

----- Switchyard headers -----> <----- Your packet header(raw bytes) ----->
<-- Payload in raw bytes -----
-----
UDP Hdr | Sequence number(32 bits) | Length(16 bits) | ETH Hdr | IP Hdr |
payload ----- Variable length
-----

```

`sequence_number` 自然就是序列号, `length` 不用多说是 在 `blaster` 中统一的, 剩下的数据部分随意发挥。

再设置一下收发地址就行了, 目的ip是 `blastee` 的ip, 目的mac我选择的是 `middlebox` 与之相连的端口的mac地址

`process_pkt()` 函数的实现如下:

```
def process_pkt(self, pkt, seq):
    pkt[0].src = '10:00:00:00:00:01'
    pkt[0].dst = '40:00:00:00:00:01'
    pkt[1].src = '192.168.100.1'
    pkt[1].dst = '192.168.200.1'
    pkt[1].ttl = 10
    data = seq.to_bytes(4, 'big')
    data += self.length.to_bytes(2, 'big')
    pkt += data
    payload = b'data ddata dddata ddddata ddddata'
    payload = payload[0:self.length-1]
    pkt += payload
```

处理完超时情况, 现在需要对 `rhs` 做出处理, 不可能只关注 `lhs` 吧, `rhs` 也得往前推进。

如果 `rhs - lhs + 1 <= sw` 并且 `rhs <= num` 时, 这项工作才有可能进行。

此时`rhs`、`lhs`都是合法的, 那检查`rhs`是否发过去包了, 如果没发, 那就不用 `process_pkt()` 包装一下, 发出去, 并且如果是第一次发, 记录一下 `begin_time`, 将 `rhs` 置为已发过包的状态, 如果`rhs`和`lhs`之间的距离还能拉长, 则`rhs`往右移动。

代码如下:

```
if self.rhs - self.lhs < 5 and self.rhs <= self.num:
    if self.queue[self.rhs].is_sent == 0:
        if self.rhs == 1 and self.begin_time == -1:
            self.begin_time = time.time()
        self.process_pkt(pkt, self.rhs)
        print("send ", self.rhs)
        self.output_byte += self.length
        self.net.send_packet("blaster-eth0", pkt)
        self.queue[self.rhs].is_sent = 1
        #self.queue[self.rhs].timer = time.time()
    if self.rhs - self.lhs < 4 and self.rhs < self.num:
        self.rhs += 1
elif self.rhs - self.lhs < 4 and self.rhs < self.num:
    self.rhs += 1
```

在while循环中, 需要一个条件来结束循环, 我选择判断 `ack_num` 是否和 `num` 相等。

```

while True:
    try:
        recv = self.net.recv_packet(timeout=1.0)
    except NoPackets:
        self.handle_no_packet()
        continue
    except Shutdown:
        break

    self.handle_packet(recv)
    if self.ack_num == self.num:
        break

```

完成了 `handle_packet()` 和 `handle_no_packet()` 函数的编写，就基本上完成了 `blaster` 的构建了。

测试结果

现在来看下测试结果，我在 `middle_box` 丢包时打印了一下，在 `blaster` 收到 `ack` 包时打印，在 `blaster` 发包时打印，并且在 `blaster` 重发时打印，将 `timeout` 置为10，`num` 置为50，加快测试速度，结果如下：

这些包是在middlebox上丢掉的。

```

7 packet dropped
13 packet dropped
17 packet dropped
20 packet dropped
38 packet dropped
38 packet dropped
47 packet dropped
47 packet dropped
47 packet dropped

```

能够看到比如7、13是丢了一次后重发接收到了的

可以在 `blaster` 的打印信息中看到相应的解释：

```

Didn't receive anything
RHS: 7 LHS 7
send 7
Didn't receive anything
RHS: 8 LHS 7
send 8
ack 8
Didn't receive anything
RHS: 9 LHS 7
send 9
ack 9
Didn't receive anything
RHS: 10 LHS 7
send 10
ack 10
Didn't receive anything
RHS: 11 LHS 7
send 11
ack 11
Didn't receive anything
RHS: 11 LHS 7
Didn't receive anything
RHS: 11 LHS 7
Didn't receive anything
RHS: 11 LHS 7
7 timeout resend
ack 7
Didn't receive anything
RHS: 12 LHS 12

```


可以看到在丢失了7之后，blaster是继续之后的发包的，不过到了11就停下了，等到发现超时，就重发了7的包，可以看到打印了一行

`7 : timeout send`，就是重发了7这个包，立马收到了7的ack，然后将rhs和lhs都更新为12了。

但比如38、47是重发的包依然丢了，以38为例。

```
Didn't receive anything
RHS: 38 LHS 38
send 38
Didn't receive anything
RHS: 39 LHS 38
send 39
ack 39
Didn't receive anything
RHS: 40 LHS 38
send 40
ack 40
Didn't receive anything
RHS: 41 LHS 38
send 41
ack 41
Didn't receive anything
RHS: 42 LHS 38
send 42
ack 42
Didn't receive anything
RHS: 42 LHS 38
Didn't receive anything
RHS: 42 LHS 38
Didn't receive anything
RHS: 42 LHS 38
38 timeout resend
Didn't receive anything
RHS: 42 LHS 38
38 timeout resend
ack 38
Didn't receive anything
RHS: 43 LHS 43
send 43
ack 43
Didn't receive anything
RHS: 44 LHS 44
```

可以看到打印了两行的 `38 : timeout send`，才正常工作。

在看下最后要求打印的数据：

```
total time: 99.05
number of retx: 9
number of coarse to: 6
throughput(bps): 59.56 bps
goodput(bps): 50.48 bps
```

可以看到一共重发了9次，一共重发了6个不同的包，整个过程持续了99.05秒

`throughput` 是要略微比 `goodput` 大的，这是源于丢包的缘故。

重新开了一次测试，用wireshark监测。

blaster:

Capturing from blaster-eth0

File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help

Apply a display filter ... <Ctrl-/>

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|-------------|---------------|---------------|----------|--------|--------------|
| 1 | 0.000000000 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 2 | 0.292210145 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 3 | 1.428500570 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 4 | 1.647682810 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 5 | 2.778883002 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 6 | 3.000355542 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 7 | 4.133666687 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 8 | 4.460740934 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 9 | 5.590765160 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 10 | 5.920294763 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 11 | 7.049453225 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 12 | 8.051609866 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 13 | 8.312411222 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |

Frame 124: 82 bytes on wire (656 bits), 82 bytes captured (656 bits) on interface blaster-eth0, id 0
Ethernet II, Src: Private_00:00:01 (10:00:00:00:00:01), Dst: 40:00:00:00:00:01 (40:00:00:00:00:01)
Internet Protocol Version 4, Src: 192.168.100.1, Dst: 192.168.200.1
User Datagram Protocol, Src Port: 0, Dst Port: 0
Data (40 bytes)

```
0000 40 00 00 00 00 01 10 00 00 00 00 01 08 00 45 00  @.....
0010 00 44 00 00 00 00 0a 11 03 56 c0 a8 64 01 c0 a8  .D.....
0020 c8 01 00 00 00 00 00 30 0b a0 00 00 00 38 00 32  .....0
0030 64 61 74 61 20 64 64 61 74 61 20 64 64 64 61 74  data dda ta
0040 61 20 64 64 64 64 61 74 61 20 64 64 64 64 61  a ddddat a
0050 74 61 ta
```

blastee:

Capturing from blastee-eth0

File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help

Apply a display filter ... <Ctrl-/>

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|---------------|---------------|---------------|----------|--------|--------------|
| 130 | 182.270200946 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 131 | 183.610555123 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 132 | 183.726715280 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 133 | 185.067213282 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 134 | 185.182249984 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 135 | 186.522102196 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 136 | 186.643186337 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 137 | 191.006404233 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 138 | 191.134774775 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 139 | 192.465761521 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 140 | 192.598293416 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |
| 141 | 194.861871439 | 192.168.100.1 | 192.168.200.1 | UDP | 82 | 0 → 0 Len=40 |
| 142 | 194.997693543 | 192.168.200.1 | 192.168.100.1 | UDP | 54 | 0 → 0 Len=12 |

Frame 1: 82 bytes on wire (656 bits), 82 bytes captured (656 bits) on interface blastee-eth0, id 0
Ethernet II, Src: 40:00:00:00:00:02 (40:00:00:00:00:02), Dst: 20:00:00:00:00:01 (20:00:00:00:00:01)
Internet Protocol Version 4, Src: 192.168.100.1, Dst: 192.168.200.1
User Datagram Protocol, Src Port: 0, Dst Port: 0
Data (40 bytes)

```
0000 20 00 00 00 00 01 40 00 00 00 00 02 08 00 45 00  .....@
0010 00 44 00 00 00 00 0a 11 03 56 c0 a8 64 01 c0 a8  .D.....
0020 c8 01 00 00 00 00 00 30 0b d7 00 00 00 01 00 32  .....0
0030 64 61 74 61 20 64 64 61 74 61 20 64 64 64 61 74  data dda ta
0040 61 20 64 64 64 64 61 74 61 20 64 64 64 64 61  a ddddat a
0050 74 61 ta
```

可以从两张截图的下方观察到，blaster发出的数据，blastee收到了，工作顺利。

实验心得

通过代码复现滑动窗口协议，加深了理解。

也对python语言有了更多的掌握，尤其是main函数中 `**kwargs` 是很神奇。

对传输层的工作原理加深了理解。